AMD DEVELOPER HACKATHON: ACT III · COMMERCE TRACK

Shelf to Listing

Photo in. Ready-to-post marketplace listing out.

THE PROBLEM

Listing something for sale is slower than owning it was worth.

Every item needs a title, a category, an honest condition call, a description, and a price guess — and most people have more than one item.

10+

items in a typical declutter

garage sales, market stalls, moving house

5-10 min

per listing, written by hand

title, description, condition, price research

0

of that gets done

the stuff just sits there instead

THE SOLUTION

One photo. One model call. One listing.

A vision-language model looks straight at the photo and returns the whole listing as structured data — no OCR step, no separate pricing lookup, no manual categorizing.

```
{
  "title": "Classic Gazelle Dutch City Bike...",
  "category": "Bicycle",
  "condition": "Good",
  "description": "Classic Dutch-style city bike...",
  "suggested_price_aud": 320
}
```

HOW IT RUNS

Built on AMD, start to finish.

Vision inference runs through Fireworks' managed, OpenAI-compatible API — backed by AMD Instinct GPUs. No self-hosted server, no separate training step.

AMD Instinct GPUs | Fireworks managed inference | Python | Streamlit

1. Photo in

upload or snap a photo of the item

2. One vision call

glm-5p3-flash on AMD Instinct, kimi-k3-fast as fallback

3. Listing out

structured JSON → rendered, ready to post

WHY IT WINS

Nobody else is doing visual input.

Across the field of public submissions, the dominant pattern is text/sensor evidence → cited answer → human approval. Shelf to Listing is the only entry that starts from a photo of a physical object — a genuinely different, instantly legible demo.

Snap a photo | Get a listing | Go sell it