AMD DEVELOPER HACKATHON: ACT III · COMMERCE TRACK

Shelf to Listing

Photo in. Ready-to-post marketplace listing out.

THE PROBLEM

Listing something for sale is slower than owning it was worth.

Every item needs a title, a category, an honest condition call, a description, and a price guess — and most people have more than one item.

10+

items in a typical declutter

garage sales, market stalls, moving house

5-10 min

per listing, written by hand

title, description, condition, price research

0

of that gets done

the stuff just sits there instead

THE SOLUTION

One photo. One model call. One listing.

A vision-language model looks straight at the photo and returns the whole listing as structured data — no OCR step, no separate pricing lookup, no manual categorizing.

```
{
  "title": "Classic Gazelle Dutch City Bike...",
  "category": "Bicycle",
  "condition": "Good",
  "description": "Classic Dutch-style city bike...",
  "suggested_price_aud": 320
}
```

WHY AMD, SPECIFICALLY

One photo at a time — not a batch job.

The honest version isn't "AMD is always cheaper": published benchmarks are workload-dependent. But our workload is single-image inference, not batched — and that's exactly the regime where AMD Instinct's larger memory per chip has shown a real cost edge.

~21%

cheaper per token

RunPod's published Mixtral 8×7B benchmark, at batch size 1

192GB

memory per chip

vs 80GB on the equivalent Nvidia part

10×

items, not 10× the cost

matters once a flipper checks more than one

`AMD Instinct GPUs` `Fireworks managed inference` `Python` `Streamlit`

WHO IT'S FOR

Three use cases, one photo.

Sellers
Clearing out a pile of stuff with no time to write ten descriptions.

Flippers
Ten items at an estate sale, ten price checks before deciding what's worth buying.

Everyone else
Garage clean-out, inherited junk — don't know what it is? Point the camera and find out.

WHY IT WINS

Nobody else is doing visual input.

Across the field of public submissions, the dominant pattern is text/sensor evidence → cited answer → human approval. Shelf to Listing is the only entry that starts from a photo of a physical object — a genuinely different, instantly legible demo.

It's live, it's on GitHub, and the guitar listing in the demo took about four seconds to write.